

MLIR × Python: Building Core Compilation Pipelines in Python

Mingyang Liu, twice@apache.org

<https://github.com/PragmaTwice>

Agenda

- Introduce MLIR Python
- Write Passes in Python
- Add Rewrite Patterns in Python
- Define Dialects in Python
- Python bindings API docs

These slides are mostly example-driven!

**USE MLIR IN
YOUR COMPILER**

**MLIR PYTHON
BINDINGS FOR
IR CONSTRUCTION**

**WRITE MLIR
PASSES IN PYTHON**

**DEFINE
MLIR DIALECTS
IN PYTHON**



MLIR Python bindings

- Almost one-to-one mapping of MLIR C++ APIs, while (try to be) pythonic by design.
- Inside LLVM monorepo, DISABLE by default in CMake.
- Implemented via nanobind and MLIR C APIs (due to RTTI-related issues).
- Still experimental and incomplete.
- Used in IREE, CIRCT, Mojo, CUDA ..

MLIR Python bindings: IR Construction



```
from mlir.ir import *
from mlir.dialects import func, arith

with Context(), Location.unknown():
    f32 = F32Type.get()

    module = Module.create()
    with InsertionPoint(module.body):

        @func.func(f32, f32, f32)
        def fma(x, y, z):
            return arith.addf(x, arith.mulf(y, z))

print(module)
```

MLIR Python bindings: IR Construction



```
from mlir.ir import *
from mlir.dialects import func, arith

with Context(), Location.unknown():
    f32 = F32Type.get()

    module = Module.create()
    with InsertionPoint(module.body):

        @func.func(f32, f32, f32)
        def fma(x, y, z):
            return arith.addf(x, arith.mulf(y, z))

print(module)
```

MLIR Python bindings: IR Construction



```
from mlir.ir import *
from mlir.dialects import func, arith

with Context(), Location.unknown():
    f32 = F32Type.get()

    module = Module.create()
    with InsertionPoint(module.body):

        @func.func(f32, f32, f32)
        def fma(x, y, z):
            return arith.addf(x, arith.mul(y, z))

print(module)
```

```
module {
    func.func @fma(%arg0: f32, %arg1: f32,
%arg2: f32) -> f32 {
    %0 = arith.mul %arg1, %arg2 : f32
    %1 = arith.addf %arg0, %0 : f32
    return %1 : f32
}
```

MLIR Python bindings: Pass Manager



```
from mlir.passmanager import PassManager

pm = PassManager('any', module.context)
pm.add('convert-to-llvm')
pm.run(module.operation)

print(module)
```

MLIR Python bindings: Pass Manager



```
from mlir.passmanager import PassManager  
  
pm = PassManager('any', module.context)  
pm.add('convert-to-llvm')  
pm.run(module.operation)  
  
print(module)
```

```
module {  
    llvm.func @fma(%arg0: f32, %arg1: f32,  
%arg2: f32) -> f32 {  
        %0 = llvm.fmul %arg1, %arg2 : f32  
        %1 = llvm.fadd %arg0, %0 : f32  
        llvm.return %1 : f32  
    }  
}
```

MLIR Python bindings: Execution Engine



```
from mlir.execution_engine import ExecutionEngine
import ctypes

with Context():
    ee = ExecutionEngine(module)
    c_float_p = ctypes.c_float * 1
    res = c_float_p(0)
    ee.invoke("fma", c_float_p(1), c_float_p(2), c_float_p(3), res)
    assert res[0] == 7
```

Note: `llvm.emit_c_interface` is omitted.

Pass in Python

- Could we just write passes in Python side?
- Don't need to touch CMake, tablegen, C++ headers/sources, and recompile the whole MLIR.
- Prototyping a new pass in Python can be way faster.
- Via `pm.add(callable)`

Pass in Python: Example

```
● ● ●

def custom_pass(op, pass_):
    # do something with `op`
    pass

pm = PassManager("any")

pm.add(custom_pass)
pm.add("canonicalize, cse, convert-to-llvm, ...")

pm.run(module)
```

Rewrite Pattern in Python

- Python passes are great. But wait.. how to define rewrite patterns?
- The greedy pattern rewrite driver is available in Python!
- (Driver for dialect conversion? Not yet.)
- Two ways to define patterns:
 - Define patterns as Python callables via PatternRewriter and RewritePatternSet;
 - Define patterns as constructing IR via PDL.

Rewrite Pattern in Python: Example

```
● ● ●

def to_muli(op, rewriter):
    with rewriter.ip:
        assert isinstance(op, arith.AddIOp)
        new_op = arith.muli(op.lhs, op.rhs, loc=op.location)
    rewriter.replace_op(op, new_op.owner)

patterns = RewritePatternSet()
patterns.add(arith.AddIOp, to_muli)
patterns.add(...)

apply_patterns_and_fold_greedily(module, patterns.freeze())
```

Rewrite Pattern in Python: via PDL

```
● ● ●

m = Module.create()
with InsertionPoint(m.body):
    @pdl.pattern(benefit=1, sym_name="addi_to_mul")
    def pat():
        index_type = pdl.TypeOp(IndexType.get())
        operand0 = pdl.OperandOp(index_type)
        operand1 = pdl.OperandOp(index_type)
        op0 = pdl.OperationOp(
            name="arith.addi", args=[operand0, operand1], types=[index_type]
        )

    @pdl.rewrite()
    def rew():
        newOp = pdl.OperationOp(
            name="arith.muli", args=[operand0, operand1], types=[index_type]
        )
        pdl.ReplaceOp(op0, with_op=newOp)

frozen = PDLModule(m).freeze()
apply_patterns_and_fold_greedily(module, frozen)
```

Define Dialects in Python

- IRDL is available in MLIR Python.
- Define new dialects as constructing MLIR modules containing `irdl.dialect` ops.
- Load your dialects via `irdl.load_dialects(module)`

Define Dialects in Python: IRDL



```
m = Module.create()
with InsertionPoint(m.body):
    myint = irdl.dialect("myint")
    with InsertionPoint(myint.body):
        constant = irdl.operation_("constant")
        with InsertionPoint(constant.body):
            iattr = irdl.base(base_name="#builtin.integer")
            i32 = irdl.is_(TypeAttr.get(IntegerType.get_signless(32)))
            irdl.attributes_([iattr], ["value"])
            irdl.results_([i32], ["cst"], [irdl.Variadicity.single])

        add = irdl.operation_("add")
        with InsertionPoint(add.body):
            i32 = irdl.is_(TypeAttr.get(IntegerType.get_signless(32)))
            irdl.operands_()
                [i32, i32],
                ["lhs", "rhs"],
                [irdl.Variadicity.single, irdl.Variadicity.single],
            )
            irdl.results_([i32], ["res"], [irdl.Variadicity.single])

    irdl.load_dialects(m)
```

Define Dialects in Python: IRDSL?

- A more intuitive and simpler interface to define dialects and ops
- Generate OpView subclasses and builder functions automatically
- Experimental and NOT yet get merged

Define Dialects in Python: IRDSL?



```
myint = irdsl.Dialect("myint")
iattr = irdsl.BaseName("#builtin.integer")
i32 = irdsl.IsType(IntegerType.get_signless(32))

@myint.op("constant")
class ConstantOp:
    value = irdsl.Attribute(iattr)
    cst = irdsl.Result(i32)

@myint.op("add")
class AddOp:
    lhs = irdsl.Operand(i32)
    rhs = irdsl.Operand(i32)
    res = irdsl.Result(i32)

myint.load()
```

MLIR Python API docs

<https://mlir.llvm.org>

The screenshot shows the official MLIR website at <https://mlir.llvm.org>. The header includes the MLIR logo and the text "Multi-Level IR Compiler Framework". A navigation bar with links to "Community", "Debugging Tips", "FAQ", "Source", "Bugs", "Logo Assets", and "Youtube Channel". On the left, a sidebar lists "Home", "Governance", "Users of MLIR", "MLIR Related Publications", "Talks", "Deprecations & Current Refactoring", "Getting Started", and "Code Documentation". A dropdown menu from the "Source" link is open, showing "Doxygen", "GitHub", "Python Bindings", and "API docs", with "Python Bindings" circled in red. The main content area features a large heading "A Novel Intermediate Representation Overview" and a paragraph describing MLIR's purpose: "The MLIR project is a novel approach to building reusable and extensible domain specific compilers. It aims to address software fragmentation, improve compilation for heterogeneous hardware, and significantly reduce the cost of building domain specific compilers, by unifying them under a single compiler infrastructure." Below this is a section titled "Weekly Public Meeting" with the text "We host a **weekly public meeting** about MLIR and the ecosystem."

MLIR
Multi-Level IR Compiler Framework

Community › Debugging Tips FAQ Source › Bugs Logo Assets Youtube Channel

Home
Governance
Users of MLIR
MLIR Related Publications
Talks
Deprecations & Current Refactoring
Getting Started
Code Documentation

Doxygen
GitHub
Python Bindings
API docs

A Novel Intermediate Representation Overview

The MLIR project is a novel approach to building reusable and extensible domain specific compilers. It aims to address software fragmentation, improve compilation for heterogeneous hardware, and significantly reduce the cost of building domain specific compilers, by unifying them under a single compiler infrastructure.

Weekly Public Meeting

We host a **weekly public meeting** about MLIR and the ecosystem.

MLIR Python API docs

<https://mlir.llvm.org/python-bindings>

MLIR Python
bindings
documentation

Search

mlir namespace

MLIR Python bindings API documentation



Welcome to the MLIR Python bindings API documentation. This website is generated from MLIR Python source code and type stubs.

- [mlir namespace](#)
 - [mlir._mlir_libs._mlir.dialects.pdl](#)
 - [mlir._mlir_libs._mlir.dialects.quant](#)
 - [mlir._mlir_libs._mlir.dialects.transform](#)
 - [mlir._mlir_libs._mlir](#)
 - [mlir._mlir_libs._mlir.ir](#)
 - [mlir._mlir_libs._mlir.passmanager](#)
 - [mlir._mlir_libs._mlir.rewrite](#)
 - [mlir._mlir_libs._mlirExecutionEngine](#)
 - [mlir._mlir_libs._mlirPythonTestNanobind](#)
 - [mlir._mlir_libs](#)
 - [mlir.dialects._acc_ops_gen](#)
 - [mlir.dialects._affine_enum_gen](#)
 - [mlir.dialects._affine_ops_gen](#)
 - [mlir.dialects._amdgpu_enum_gen](#)
 - [mlir.dialects._amdgpu_ops_gen](#)

MLIR Python API docs

[https://mlir.llvm.org/python-bindings/...](https://mlir.llvm.org/python-bindings/)

```
print(large_elements_limit: int | None = None, large_resource_limit: int | None = None,
      enable_debug_info: bool = False, pretty_debug_info: bool = False,
      print_generic_op_form: bool = False, use_local_scope: bool = False,
      use_name_loc_as_prefix: bool = False, assume_verified: bool = False, file: object |
      None = None, binary: bool = False, skip_regions: bool = False) → None
```

Prints the assembly form of the operation to a file like object.

PARAMETERS:

- **large_elements_limit** – Whether to elide elements attributes above this number of elements. Defaults to None (no limit).
- **large_resource_limit** – Whether to elide resource attributes above this number of characters. Defaults to None (no limit). If large_elements_limit is set and this is None, the behavior will be to use large_elements_limit as large_resource_limit.
- **enable_debug_info** – Whether to print debug/location information. Defaults to False.
- **pretty_debug_info** – Whether to format debug information for easier reading by a human (warning: the result is unparseable). Defaults to False.
- **print_generic_op_form** – Whether to print the generic assembly forms of all ops. Defaults to False.

MLIR Python

- Try it now?

```
● ● ●  
$ pip install mlir-python-bindings -f "https://llvm.github.io/eudsl"  
  
$ python  
>>> from mlir import ir  
>>> with ir.Context():  
...     print(ir.Module.parse("arith.constant 10"))  
...  
module {  
    %c10_i64 = arith.constant 10 : i64  
}
```

MLIR Python

- Try it in your browser? Without installation!
- <https://llvm.github.io/eudsl/console> - Pyodide REPL with MLIR Python bindings loaded
- <https://llvm.github.io/eudsl/jupyter> - JupyterLab packed with MLIR Python bindings
- via the Wasm version of MLIR python.

Thanks